

Multiple Linear Regression in R

The following is a multi-linear regression assessment of a given 1985 auto imports database developed by Jeffrey C. Schlimmer and maintained at the University of California, School of Information and Computer Science data repository. The main goal during this exercise is to explore this dataset, develop a multi-linear regression model using R, and identify unique characteristics of this machine learning modeling algorithm.

This multivariate dataset was consolidated from various sources including the 1985 Ward's Automotive Yearbook, Personal Auto Manuals and Insurance Services Office, New York, and Insurance Collision Report – Insurance Institute for Highway Safety of Washington, D.C. a variety of attributes including, continuous and nominal values as well as missing values. The dataset consists of 26 attributes and 205 observations with a mix of categorical, integers, numeric and factor types of variables including missing values (Dua, D. and Graff, C. (2019).

Per the original repository, the below list describes each of the attributes of this dataset.

- symboling: -3, -2, -1, 0, 1, 2, 3.
- normalized-losses: continuous from 65 to 256.
- make: alfa-romero, audi, bmw, chevrolet, dodge, honda, isuzu, jaguar, mazda, mercedes-benz, mercury, mitsubishi, nissan, peugot, plymouth, porsche, renault, saab, subaru, toyota, volkswagen, volvo
- fuel-type: diesel, gas.
- aspiration: std, turbo.
- num-of-doors: four, two.
- body-style: hardtop, wagon, sedan, hatchback, convertible.
- drive-wheels: 4wd, fwd, rwd.
- engine-location: front, rear.
- wheel-base: continuous from 86.6 to 120.9.
- length: continuous from 141.1 to 208.1.
- width: continuous from 60.3 to 72.3.
- height: continuous from 47.8 to 59.8.
- curb-weight: continuous from 1488 to 4066.
- engine-type: dohc, dohcvt, l, ohc, ohcvt, ohcvt, rotor.
- num-of-cylinders: eight, five, four, six, three, twelve, two.
- engine-size: continuous from 61 to 326.
- fuel-system: 1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi.
- bore: continuous from 2.54 to 3.94.
- stroke: continuous from 2.07 to 4.17.
- compression-ratio: continuous from 7 to 23.
- horsepower: continuous from 48 to 288.
- peak-rpm: continuous from 4150 to 6600.
- city-mpg: continuous from 13 to 49.
- highway-mpg: continuous from 16 to 54.
- **price**: continuous from 5118 to 45400.

The dependent variable (response) of this dataset is the price of the vehicle, while the other remaining 25 variables will be considered independents or predictors. My goal is to illustrate how a multi-linear regression analysis can enable or empower auto sales and insurance services with immediate quotes and/or valuations of new or used vehicle based on some particular characteristics of the car.

Preview of the data set. The first step was to set and load the data to be analyzed. The `imports_85.csv` dataframe was loaded into the console using the `read.csv()` command and saved as a vector named `import`. The image below shows a preview of the set.

```
> #1 Set the working directory/load the data
> setwd("~/OneDrive/UMGC/DATA630/Week4/")
> # Read the csv file
> import <- read.csv("imports_85.csv", header=TRUE, sep = ",", as.is=FALSE)
> head(import)           #Preview the dataframe
```

```
> head(import)
  symboling normalized_losses fuel_type aspiration num_doors body_style drive_wheels wheel_base length width height curb_weight
4         2             164      gas      std      four   sedan         fwd       99.8   176.6   66.2   54.3       2337
5         2             164      gas      std      four   sedan         4wd       99.4   176.6   66.4   54.3       2824
7         1             158      gas      std      four   sedan         fwd      105.8   192.7   71.4   55.7       2844
9         1             158      gas     turbo      four   sedan         fwd      105.8   192.7   71.4   55.9       3086
11        2             192      gas      std      two    sedan         rwd      101.2   176.8   64.8   54.3       2395
12        0             192      gas      std      four   sedan         rwd      101.2   176.8   64.8   54.3       2395
  engine_size bore stroke compression_ratio horsepower peak_rpm city_mpg highway_mpg Price
4         109  3.19   3.4           10.0         102    5500      24         30  13950
5         136  3.19   3.4            8.0         115    5500      18         22  17450
7         136  3.19   3.4            8.5         110    5500      19         25  17710
9         131  3.13   3.4            8.3         140    5500      17         20  23875
11        108  3.50   2.8            8.8         101    5800      23         29  16430
12        108  3.50   2.8            8.8         101    5800      23         29  16925
```

Data preprocessing. During the data preprocessing phase, and as directed, I excluded four of the original dataframe (df) variables using the `NULL` command. Those variables were `engine_type`, `make`, `num_of_cylinders`, and `fuel_system`.

```
> #2 Data Pre Processing
> import$engine_type<-NULL           #Remove the variable
> import$make<-NULL                 #Remove the variable
> import$num_of_cylinders<-NULL      #Remove the variable
> import$fuel_system<-NULL           #Remove the variable
```

Following this action, I checked the data structure with `str()` and noticed NA values across some of the attributes (figure 1.1). After familiarizing myself with these missing values, I decided to ignore these observations containing NA values with the `na.omit()` command. In retrospect, this data preparation step proved to be critical in preparation of the modeling of the data. The newly updated df was reassigned to the “import” dataset. Then, to confirm my previous actions, and confirmed the content of the df, I pulled a `summary()` command as shown in figure 1.2.

```
str(import) #Check data structure
```

```
> str(import) #Check data structure
'data.frame': 205 obs. of 22 variables:
 $ symboling : int 3 3 1 2 2 2 1 1 1 0 ...
 $ normalized_losses: int NA NA NA 164 164 NA 158 NA 158 NA ...
 $ fuel_type : Factor w/ 2 levels "diesel","gas": 2 2 2 2 2 2 2 2 2 ...
 $ aspiration : Factor w/ 2 levels "std","turbo": 1 1 1 1 1 1 1 1 2 ...
 $ num_doors : Factor w/ 2 levels "four","two": 2 2 2 1 1 2 1 1 1 2 ...
 $ body_style : Factor w/ 5 levels "convertible",...: 1 1 3 4 4 4 4 5 4 3 ...
 $ drive_wheels : Factor w/ 3 levels "4wd","fwd","rwd": 3 3 3 2 1 2 2 2 2 1 ...
 $ engine_location : Factor w/ 2 levels "front","rear": 1 1 1 1 1 1 1 1 1 1 ...
 $ wheel_base : num 88.6 88.6 94.5 99.8 99.4 ...
 $ length : num 169 169 171 177 177 ...
 $ width : num 64.1 64.1 65.5 66.2 66.4 66.3 71.4 71.4 71.4 67.9 ...
 $ height : num 48.8 48.8 52.4 54.3 54.3 53.1 55.7 55.7 55.9 52 ...
 $ curb_weight : int 2548 2548 2823 2337 2824 2507 2844 2954 3086 3053 ...
 $ engine_size : int 130 130 152 109 136 136 136 136 131 131 ...
 $ bore : num 3.47 3.47 2.68 3.19 3.19 3.19 3.19 3.19 3.13 3.13 ...
 $ stroke : num 2.68 2.68 3.47 3.4 3.4 3.4 3.4 3.4 3.4 3.4 ...
 $ compression_ratio: num 9 9 9 10 8 8.5 8.5 8.5 8.3 7 ...
 $ horsepower : int 111 111 154 102 115 110 110 110 140 160 ...
 $ peak_rpm : int 5000 5000 5000 5500 5500 5500 5500 5500 5500 5500 ...
 $ city_mpg : int 21 21 19 24 18 19 19 19 17 16 ...
 $ highway_mpg : int 27 27 26 30 22 25 25 25 20 22 ...
 $ Price : int 13495 16500 16500 13950 17450 15250 17710 18920 23875 23870 ...
> |
```

Figure 1.1 – The import data frame contains 205 observations across 22 variables.

```
import<-na.omit(import) #Exclude rows with NA values
summary(import) #Check the descriptive statistics
```

```
> import<-na.omit(import) #Exclude variables w/ NA values in prep. for model
> summary(import) #Check the descriptive statistics
```

symboling	normalized_losses	fuel_type	aspiration	num_doors	body_style	drive_wheels	engine_location
Min. :-2.0000	Min. : 65.0	diesel: 15	std :132	four:96	convertible: 2	4wd: 8	front:160
1st Qu.: 0.0000	1st Qu.: 94.0	gas :145	turbo: 28	two :64	hardtop : 5	fwd:106	rear : 0
Median : 1.0000	Median :114.0				hatchback :56	rwd: 46	
Mean : 0.7375	Mean :121.3				sedan :80		
3rd Qu.: 2.0000	3rd Qu.:148.0				wagon :17		
Max. : 3.0000	Max. :256.0						

wheel_base	length	width	height	curb_weight	engine_size	bore	stroke
Min. : 86.60	Min. :141.1	Min. :60.3	Min. :49.40	Min. :1488	Min. : 61.0	Min. :2.540	Min. :2.070
1st Qu.: 94.50	1st Qu.:165.5	1st Qu.:64.0	1st Qu.:52.00	1st Qu.:2073	1st Qu.: 97.0	1st Qu.:3.050	1st Qu.:3.107
Median : 96.90	Median :172.2	Median :65.4	Median :54.10	Median :2338	Median :110.0	Median :3.270	Median :3.270
Mean : 98.24	Mean :172.3	Mean :65.6	Mean :53.88	Mean :2459	Mean :119.1	Mean :3.298	Mean :3.237
3rd Qu.:100.60	3rd Qu.:177.8	3rd Qu.:66.5	3rd Qu.:55.50	3rd Qu.:2809	3rd Qu.:134.5	3rd Qu.:3.550	3rd Qu.:3.410
Max. :115.60	Max. :202.6	Max. :71.7	Max. :59.80	Max. :4066	Max. :258.0	Max. :3.940	Max. :4.170

compression_ratio	horsepower	peak_rpm	city_mpg	highway_mpg	Price
Min. : 7.00	Min. : 48.00	Min. :4150	Min. :15.00	Min. :18.00	Min. : 5118
1st Qu.: 8.70	1st Qu.: 69.00	1st Qu.:4800	1st Qu.:23.00	1st Qu.:28.00	1st Qu.: 7384
Median : 9.00	Median : 88.00	Median :5200	Median :26.00	Median :32.00	Median : 9164
Mean :10.15	Mean : 95.88	Mean :5116	Mean :26.51	Mean :32.07	Mean :11428
3rd Qu.: 9.40	3rd Qu.:114.00	3rd Qu.:5500	3rd Qu.:31.00	3rd Qu.:37.00	3rd Qu.:14559
Max. :23.00	Max. :200.00	Max. :6600	Max. :49.00	Max. :54.00	Max. :35056

```
> |
```

Figure 1.2 – Descriptive statistics of the dataset. The engine_location variable appears as a one-value variable.

After omitting the NA values and running the summary command, I noticed the engine_location variable with only one-type of value (front = 160); however, I decided to press forward without any additional changes. Further in the process, as I was building the model, RStudio gave me the below “error in contrasts”, which urged me to return to this step and remove the one-value variable of engine_location using the NULL command.

```
model<-lm(Price~., train.data)
```

```
Error in `contrasts<-`(`*tmp*`, value = contr.funs[1 +
isOF[nn]]):contrasts can be applied only to factors with 2 or more
levels
```

```
> import$engine_location<-NULL # Remove the one-value variable>
summary(import) #Check the descriptive statistics
```

```
> import$engine_location<-NULL # Remove one-value variable in prep. for model
> summary(import) #Check the descriptive statistics
```

symboling	normalized_losses	fuel_type	aspiration	num_doors	body_style	drive_wheels	wheel_base
Min. : -2.0000	Min. : 65.0	diesel: 15	std :132	four:96	convertible: 2	4wd: 8	Min. : 86.60
1st Qu.: 0.0000	1st Qu.: 94.0	gas :145	turbo: 28	two :64	hardtop : 5	fwd:106	1st Qu.: 94.50
Median : 1.0000	Median :114.0				hatchback :56	rwd: 46	Median : 96.90
Mean : 0.7375	Mean :121.3				sedan :80		Mean : 98.24
3rd Qu.: 2.0000	3rd Qu.:148.0				wagon :17		3rd Qu.:100.60
Max. : 3.0000	Max. :256.0						Max. :115.60

length	width	height	curb_weight	engine_size	bore	stroke	compression_ratio
Min. :141.1	Min. :60.3	Min. :49.40	Min. :1488	Min. : 61.0	Min. :2.540	Min. :2.070	Min. : 7.00
1st Qu.:165.5	1st Qu.:64.0	1st Qu.:52.00	1st Qu.:2073	1st Qu.: 97.0	1st Qu.:3.050	1st Qu.:3.107	1st Qu.: 8.70
Median :172.2	Median :65.4	Median :54.10	Median :2338	Median :110.0	Median :3.270	Median :3.270	Median : 9.00
Mean :172.3	Mean :65.6	Mean :53.88	Mean :2459	Mean :119.1	Mean :3.298	Mean :3.237	Mean :10.15
3rd Qu.:177.8	3rd Qu.:66.5	3rd Qu.:55.50	3rd Qu.:2809	3rd Qu.:134.5	3rd Qu.:3.550	3rd Qu.:3.410	3rd Qu.: 9.40
Max. :202.6	Max. :71.7	Max. :59.80	Max. :4066	Max. :258.0	Max. :3.940	Max. :4.170	Max. :23.00

horsepower	peak_rpm	city_mpg	highway_mpg	Price
Min. : 48.00	Min. :4150	Min. :15.00	Min. :18.00	Min. : 5118
1st Qu.: 69.00	1st Qu.:4800	1st Qu.:23.00	1st Qu.:28.00	1st Qu.: 7384
Median : 88.00	Median :5200	Median :26.00	Median :32.00	Median : 9164
Mean : 95.88	Mean :5116	Mean :26.51	Mean :32.07	Mean :11428
3rd Qu.:114.00	3rd Qu.:5500	3rd Qu.:31.00	3rd Qu.:37.00	3rd Qu.:14559
Max. :200.00	Max. :6600	Max. :49.00	Max. :54.00	Max. :35056

Figure 1.3 – Summary of the dataset after completion of preprocessing.

At this point, I was comfortable with the structure and content of the dataset and advanced to the transformation phase of splitting the original dataset into a training and test data subsets.

Training and test data. The first step in creating the subsets was establishing a seed. The `set.seed` command was used to ensure the results were reproducible. More specifically, the created seed ensures that the data would get divided similarly every time (Bansal, 2020).

In this model, the training set was set to hold 70% of the observations, while the test data, the remaining 30%. The main reason we split the dataset between training and data is for evaluating the performance and accuracy of the generated algorithm. This technique is vital when handling large data sets, or when developing a time-intensive model, and/or accuracy or specific threshold is expected of the model performance (Brownlee, 2020). Below (see figure 1.4A, & figure 1.4B,) are the used commands to set the seed and divide the data into the training and test subsets, followed by a quick glimpse of each subset by the `head()` command.

```
# Set the seed value to ensure that result were reproducible
set.seed(1234)
> # Divide the data into training and test data
```

```
> smpl <- sample(2, nrow(import), replace = TRUE, prob = c(0.7, 0.3))
> train.data <- import [smpl == 1, ]      #Training sample
> head(train.data)
```

```
> train.data <- import [smpl == 1, ]      #Training sample
> head(train.data)
  symboling normalized_losses fuel_type aspiration num_doors body_style drive_wheels engine_location wheel_base length width
4         2             164      gas      std      four      sedan         fwd         front       99.8  176.6  66.2
5         2             164      gas      std      four      sedan         4wd         front       99.4  176.6  66.4
7         1             158      gas      std      four      sedan         fwd         front      105.8  192.7  71.4
9         1             158      gas     turbo      four      sedan         fwd         front      105.8  192.7  71.4
12        0             192      gas      std      four      sedan         rwd         front      101.2  176.8  64.8
13        0             188      gas      std      two       sedan         rwd         front      101.2  176.8  64.8
  height curb_weight engine_size bore stroke compression_ratio horsepower peak_rpm city_mpg highway_mpg Price
4   54.3      2337      109 3.19  3.40           10.0         102      5500      24       30  13950
5   54.3      2824      136 3.19  3.40           8.0         115      5500      18       22  17450
7   55.7      2844      136 3.19  3.40           8.5         110      5500      19       25  17710
9   55.9      3086      131 3.13  3.40           8.3         140      5500      17       20  23875
12  54.3      2395      108 3.50  2.80           8.8         101      5800      23       29  16925
13  54.3      2710      164 3.31  3.19           9.0         121      4250      21       28  20970
>
```

Figure 1.4A – Heading of training dataset.





```
> test.data <- import [smpl == 2, ]      #Test sample
> head(test.data)
```

```
> test.data <- imp [ind == 2, ]
> head(test.data)
  symboling normalized_losses fuel_type aspiration num_doors body_style drive_wheels wheel_base length width height curb_weight
11        E             192      gas      std      two       sedan         rwd       101.2  176.8  64.8  54.3      2395
24        D             118      gas     turbo      two  hatchback         fwd       93.7  157.3  63.8  50.8      2128
26        D             148      gas      std      four      sedan         fwd       93.7  157.3  63.8  50.6      1989
36        C             110      gas      std      four      sedan         fwd       96.5  163.4  64.0  54.5      2010
38        C             106      gas      std      two  hatchback         fwd       96.5  167.5  65.2  53.3      2236
39        C             106      gas      std      two  hatchback         fwd       96.5  167.5  65.2  53.3      2289
  engine_size bore stroke compression_ratio horsepower peak_rpm city_mpg highway_mpg Price
11  108 (3,3.5] (2.49,2.91]           8.8         101      5800      23       29  16430
24   98 (3,3.5] (3.33,3.75]           7.6         102      5500      24       30   7957
26   90 (2.5,3] (2.91,3.33]           9.4          68      5500      31       38   6692
36   92 (2.5,3] (3.33,3.75]           9.2          76      6000      30       34   7295
38  110 (3,3.5] (3.33,3.75]           9.0          86      5800      27       33   7895
39  110 (3,3.5] (3.33,3.75]           9.0          86      5800      27       33   9095
>
```


Figure 1.4B – Heading of testing dataset .

After splitting the data, I continued to create the model. I used the `lm()` command to fit the regression model with the training data. Here, price was set as the dependable variable and the rest of the variables – using the any-character wildcard (i.e. `Price ~ ., ...`)—as predictors.


```
# Build the model w/ lm command.
> model<-lm(Price~., train.data)
> plot(model, col = "blue")
```




The `plot()` command generated four unique plots of the model. The first one displayed the model's residuals vs. fitted values. As observed, the model reflects positive residuals (above the dotted line), and negative residuals (below the dotted line). The curved red line highlights the pattern followed by the error residuals.



The second graphic shows a normal q-q plot displaying the distribution of the residuals against the quartiles. Per this graphic we can argue that the majority of the residuals are normally distributed as they run close to the dotted line. Some exceptions, observations 66, 109, and 120.



Next, the scale-location plot. This plot indicates how spread were the points along the predicted range. Per Shantanu Deo, one of the assumptions for regression is that the points' variance should be within a reasonable range of the predictor (2016). In this model, we may say that between 5000 and 20000 the variance shows a good degree of uniformity.



Finally, the residuals and leverage plot. This graphic displays how much leverage, and therefore, influence a point may exerts onto the regression model's variable of price, should a given observation be removed. In terms of the Cook's distance red dotted-line reference, any point falling above or below this line are considered of high leverage. In this case, no point was found within the mentioned regions.




Figure 1.5 – Model's residual plots .

The next step was reviewing the model's descriptive statistics using the `summary()` command. Displayed below, this command produced particulars of the model including the formula used, residuals information, and coefficients.

```
> summary(model)      # Run summary( ) to see descriptive statistics
Call:
lm(formula = Price ~ ., data = train.data)

Residuals:
    Min       1Q   Median       3Q      Max
-6273.9 -1238.0  -106.9   1045.8   5942.4

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   -6.447e+04  2.177e+04  -2.962  0.00387 **
fuel_typegas    1.773e+04  8.003e+03   2.215  0.02919 *
aspirationturbo  3.180e+03  1.162e+03   2.736  0.00744 **
wheel_base     2.279e+02  1.338e+02   1.704  0.09176 .
curb_weight     5.963e+00  3.115e+00   1.914  0.05861 .
bore          -3.647e+03  1.320e+03  -2.764  0.00688 **
compression_ratio 1.292e+03  5.847e+02   2.210  0.02950 *
---
... "See appendix for complete summary\( \) results"

Residual standard error: 2294 on 94 degrees of freedom
Multiple R-squared:  0.8791,    Adjusted R-squared:  0.8482
F-statistic: 28.48 on 24 and 94 DF,  p-value: < 2.2e-16
```

Observations about the summary output

How does the model represent the relationships between dependent and independent variables in the auto import dataset?

The model represents the relationship between the predictors and the dependent variables. Each coefficients' estimated values underline the relationship between such variable and the targeted variable. The residuals range gives an example of how normally distributed the residuals are, important aspect of regression. Other values such as the adjusted r-squared, p-value and F-statistics give additional parameters to confirm the veracity of the model.

How does the method handle categorical variables?

Categorical variables are handled differently. Case in point, across the df the aspiration variable only shows two options, std or turbo. The model evaluates one of the two, in this case aspiration_turbo. Basically, when the value is zero this coefficient does not influence the

total Price, but when it is equal to one (meaning the car has turbo), the price goes up by \$3,180

What does the residuals section of the output mean?

The residuals section refers to the observed values of the error term for each of the given observations. In simple terms, residuals are the opposite side of the predictions, the distances between the observed and predicted values. In this model, the residuals' variation runs from -6273 to 5942 with a median around -100. As previously mentioned, these noted range and median illustrates the models' quasi-normal distribution and constant variance (Dietrich, Heller, & Yang, 2015).

What are coefficients, and what do they mean?

The coefficient is an estimated calculation for each variable, based on the ordinary least squares (OLS), which aims to minimize the error delta between the linear model and the actual observations, in order to estimate and trace the fit line of the model that best approximates the relationship between the outcome variable (Price) and the independent variables (Dietrich, Heller, & Yang, 2015). Consider the curb_weight value of 5.963. It means that for every single unit increase in a car's curb_weight the total price increases by \$5,963.

What is an intercept, and what does it mean?

The intercept refers to the value of y when x = 0. When everything still at zero, the point where the line touches the vertical or y axis, denoted in the below multiple linear regression equation as β_0 (Triola, 2017).

$$y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_3x_3 + \cdots + \beta_{p-1}x_{p-1} + \varepsilon$$

What do the p-values tell us about the significance of each variable?

The p-values indicates how statistically significant a variable is in relation to either supporting or rejecting a null hypothesis. Normally a p-value less than 0.05 or 5% will be considered as a strong evidence to reject the null hypothesis, and support the other position, in this case the Price as the response.

What is the overall accuracy of the model?

The accuracy of the model can be measured in terms of the model's residual standard error (RSE). In other words, the RSE quantifies that threshold between the observations and the predicted regression line, thus describing the accuracy of the model. The lower the

number the better accuracy the model has. For this model the RSE is 2294 on 94 degrees of freedom.

Evaluate the model on a test set. Once completed with the training data model, I moved to test the model with the 30% train data set. For this step, I used the `predict()` command which take in consideration the `test.data` argument and evaluates these new set of variables against the model created with the `lm()` command. Statistically, the prediction model displayed a minimum value of 5255, a max of 26997 and a median around 8850.

```
# Evaluate the model on test data
pred <- predict(model, newdata=test.data)
> summary(pred)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  5255    7432    8858   10585   12742   26997
```

Then I moved to display the results of predicted vs. observed values using the `plot()` command. Here I use the vectored prediction variable of `pred` to assess the `test.data` subset. Once generated, I added a fit line to compare actual vs. predicted values.

```
# Evaluate the model on test data
pred <- predict(model, newdata=test.data)
# Test data scatter plot
plot(test.data$Price, pred, xlab = "Observed", ylab = "Prediction",
     main = "Model Evaluation on Test Data", col = "dark red")
abline(a = 0, b = 1, col = "blue")
```

Importantly, this visual representation of the test data, confirms some of the details mentioned about the underpinning coefficients of low p-value and coefficient of determination of 84%. Per the graphic, the majority of the observations were along the prediction line, confirming that the residuals had a good distribution and minimal variance.




Figure 1.6 – Test set conforms to the model's prediction.

The close distance between the points and the fit line, measured by the ordinary least squares, talks about the small overall error and usefulness of the developed model.

As previously mentioned on the residuals plot, I found the scale-location plot highly informative. The plot illustrates the residuals' variance along the prediction line. In addition, I created a histogram (figure 1.7) to better appreciate the residuals distribution, which once again confirmed the aforementioned values of an almost symmetrical distribution.

```
> model<-lm(Price~., train.data)
> par(mfrow = c(1,1))
> plot(model, col = "blue")
Hit <Return> to see next plot:
# Additional histogram to show residuals distribution
> hist(model$residual, col = blues9)
```




Figure 1.7 – Residuals distribution well-balanced.

Minimal adequate model. The minimal adequate model (MAM) involves the minimum set of variables or predictors needed to sustain the model at hand. This is a minimalist approach that follows the principle of parsimony (aka Occam's razor) which states that a model should be as simple as possible (Bruce, P & Bruce, A. 2017). This approach can result very useful when processing big datasets with numerous attributes and observations, by concentrating performance and processing efforts into particular variables with strong leverage.

In building a reduced model I used the `step()` command with a backward direction. This command generated many different steps, each one with different AIC values. The Akaike's Information Criteria (AIC) is a metric that measures the level of observations and variables used for the model. Following the Occam's razor principle, the goal is to identify a model with the

lowest AIC score (Bruce, P & Bruce, A. 2017). The below command output shows the step with the lowest AIC equal to 1839.57, and only using the normalized_losses, fuel_type, aspiration, wheel_base, curb_weight, engine_size, bore, and compression_ratio.

```
# Use the step function to build a reduced model
model2<-step(model, direction="backward")
# Preview model2 and identify lowest-value AIC
Step:  AIC=1839.57
Price ~ normalized_losses + fuel_type + aspiration + wheel_base +
      curb_weight + engine_size + bore + compression_ratio
```

	Df	Sum of Sq	RSS	AIC
<none>			528983234	1839.6
- normalized_losses	1	15431693	544414927	1841.0
- engine_size	1	19872679	548855913	1842.0
- wheel_base	1	30890707	559873941	1844.3
- bore	1	45531578	574514812	1847.4
- compression_ratio	1	55921299	584904533	1849.5
- fuel_type	1	57462282	586445516	1849.8
- aspiration	1	83448857	612432091	1855.0
- curb_weight	1	108234656	637217890	1859.7

To further analyze this step, I used the summary () command for this second model.

```
> summary(model2)
```

```
> summary(model2)

Call:
lm(formula = Price ~ normalized_losses + fuel_type + aspiration +
    wheel_base + curb_weight + engine_size + bore + compression_ratio,
    data = train.data)

Residuals:
    Min       1Q   Median       3Q      Max
-6447.9 -1090.4  -51.5    920.3   6502.8

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  -57061.666   12507.494  -4.562 1.32e-05 ***
normalized_losses    11.719     6.542   1.791 0.075986 .
fuel_typegas    20628.663   5967.663   3.457 0.000778 ***
aspirationturbo  3380.538    811.522   4.166 6.20e-05 ***
wheel_base      201.644     79.560   2.534 0.012668 *
curb_weight       7.648     1.612   4.744 6.33e-06 ***
engine_size      36.753     18.079   2.033 0.044477 *
bore           -3137.874    1019.774  -3.077 0.002639 **
compression_ratio 1488.471     436.492   3.410 0.000909 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2193 on 110 degrees of freedom
Multiple R-squared:  0.8707,    Adjusted R-squared:  0.8613
F-statistic: 92.6 on 8 and 110 DF,  p-value: < 2.2e-16
```

```
> model2
```

```
Call:
lm(formula = Price ~ normalized_losses + fuel_type + aspiration +
    wheel_base + curb_weight + engine_size + bore + compression_ratio,
    data = train.data)

Coefficients:
(Intercept)  normalized_losses      fuel_typegas  aspirationturbo      wheel_base      curb_weight
   -57061.666         11.719        20628.663         3380.538         201.644           7.648
   engine_size           bore  compression_ratio
     36.753       -3137.874         1488.471
```

What are the coefficients and the intercept, and what do they mean?

For the second model, the only displayed coefficients were the 8 identified ones during the `step()` command and criteria of minimum AIC.

Compare the prediction accuracy of the minimum adequate model with the prediction accuracy of the original model.

First, to compare both models, the original one and the MAM, I used the `AIC()` command as illustrated below. The second model, had the lowest AIC across 10 degrees of freedom, while the original training dataset, which included all the independent variables was scored higher than the MAM. This number confirms the fact that more variables does not necessarily means better accuracy of a model.

```
> AIC(model, model2) # Use AIC function to evaluate both models
      df      AIC
model  26 2203.309
model2 10 2179.281
```

Secondly, I compared both models statistic values as shown below. Notice the lower residual standard error (RSE), which relates to the accuracy of the models, between model 2 and model 1. Additionally, the adjusted r-squared, highlights the ratio of included observations part of the model increases from 84% to 86%. Lastly, the F-statistic values also increases from 28.4 to 92.6.

```
Residual standard error: 2294 on 94 degrees of freedom
Multiple R-squared:  0.8791, Adjusted R-squared:  0.8482      # Model 1 RSE
F-statistic: 28.48 on 24 and 94 DF,  p-value: < 2.2e-16

Residual standard error: 2193 on 110 degrees of freedom
Multiple R-squared:  0.8707, Adjusted R-squared:  0.8613      # Model 2 RSE
F-statistic:  92.6 on 8 and 110 DF,  p-value: < 2.2e-16
```

Suppose that we have a new car, and we know the values for the independent variables. How would you use the model to predict the value of the dependent variable for the new car?

The model gives you a baseline to anchor your deductions. If we have the value of those 8 independent variables identified as part of the step with the lowest AIC, we will be able to estimate the total price of the vehicle by substituting or comparing the new parameters against the model, all other factors remaining equal.

In conclusion, this exercise aimed to explain how a multi-linear regression model could result in beneficial predicting or anticipating continuous values based on the correlation between a targeted variable and its correlated predictors. The exercise provides the opportunity to utilize different commands with R to assess a dataset, build a regression model, and test against a subset of the original dataset. As it relates to the auto dataset, in fact the model was able to predict, with a high level of confidence, the anticipated values in price. Along the assessment, it is noticed that the accuracy of the model could be impacted by the residuals' distribution of the tested population or sampled data. In other words, the model may not be as accurate assessing asymmetrical distributions.

References

- Bansal, J. (2020). How to Use Random Seeds Effectively. <https://towardsdatascience.com/how-to-use-random-seeds-effectively-54a4cd855a79>
- Brownlee, J. (2020). Train-test Split for Evaluating Machine Learning Algorithms. <https://machinelearningmastery.com/train-test-split-for-evaluating-machine-learning-algorithms/>
- Deo, S. (2016). R Tutorial : Residual Analysis for Regression. <http://analyticspro.org/2016/03/05/r-tutorial-residual-analysis-for-regression/>
- Dietrich, D., Heller, B., & B. Yang. (2015). Data science & big data analytics: Discovering, analyzing, visualizing and presenting data. John Wiley & Sons, Inc.
- Dua, D. and Graff, C. (2019). UCI Machine Learning Repository [<http://archive.ics.uci.edu/ml>]. Irvine, CA: University of California, School of Information and Computer Science.
- Bruce, B., & Bruce, A. (2017). Practical statistics for data science. O'Reilly Media, Inc.

Appendix

```
> summary(model)
```

Call:
lm(formula = Price ~ ., data = train.data)

Residuals:

Min	1Q	Median	3Q	Max
-4766.4	-1244.5	-130.1	774.3	5443.3

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-8.553e+04	2.195e+04	-3.896	0.00019	***
symboling	-2.185e+02	3.514e+02	-0.622	0.53566	
normalized_losses	1.546e+01	9.721e+00	1.590	0.11534	
fuel_typegas	1.068e+04	8.225e+03	1.299	0.19738	
aspirationturbo	3.175e+03	1.282e+03	2.476	0.01519	*
num_doorstwo	2.270e+02	8.070e+02	0.281	0.77916	
body_stylehatchback	6.605e+02	1.276e+03	0.517	0.60611	
body_stylededan	1.098e+03	1.366e+03	0.803	0.42391	
body_stylewagon	1.284e+03	1.617e+03	0.794	0.42914	
drive_wheelsfwd	-6.753e+02	1.287e+03	-0.525	0.60107	
drive_wheelsrwd	1.050e+03	1.558e+03	0.674	0.50231	
wheel_base	1.068e+02	1.528e+02	0.699	0.48643	
length	-5.721e+01	6.571e+01	-0.871	0.38625	
width	8.691e+02	2.837e+02	3.064	0.00290	**
height	2.339e+02	1.956e+02	1.196	0.23494	
curb_weight	1.193e+00	2.228e+00	0.535	0.59366	
engine_size	6.308e+01	3.213e+01	1.963	0.05277	.
bore	-1.377e+03	1.291e+03	-1.067	0.28902	
stroke	-1.478e+03	9.386e+02	-1.575	0.11892	
compression_ratio	9.262e+02	5.905e+02	1.568	0.12040	
horsepower	9.403e+00	2.910e+01	0.323	0.74739	
peak_rpm	7.211e-01	7.863e-01	0.917	0.36156	
city_mpg	-3.511e+02	2.086e+02	-1.683	0.09589	.
highway_mpg	2.505e+02	1.903e+02	1.317	0.19137	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2282 on 88 degrees of freedom
Multiple R-squared: 0.8664, Adjusted R-squared: 0.8314
F-statistic: 24.8 on 23 and 88 DF, p-value: < 2.2e-16

```
> coef(model) # To display the model coefficients only
```

(Intercept)	symboling	normalized_losses	fuel_typegas	aspirationturbo	num_doorstwo
-8.553431e+04	-2.185308e+02	1.546071e+01	1.068288e+04	3.174741e+03	2.269953e+02
body_stylehatchback	body_stylededan	body_stylewagon	drive_wheelsfwd	drive_wheelsrwd	wheel_base
6.605023e+02	1.097810e+03	1.284263e+03	-6.753284e+02	1.049831e+03	1.067810e+02
length	width	height	curb_weight	engine_size	bore
-5.721460e+01	8.690945e+02	2.338774e+02	1.192926e+00	6.308358e+01	-1.376596e+03
stroke	compression_ratio	horsepower	peak_rpm	city_mpg	highway_mpg
-1.477951e+03	9.261566e+02	9.402638e+00	7.211378e-01	-3.510797e+02	2.505175e+02

[illegible]